

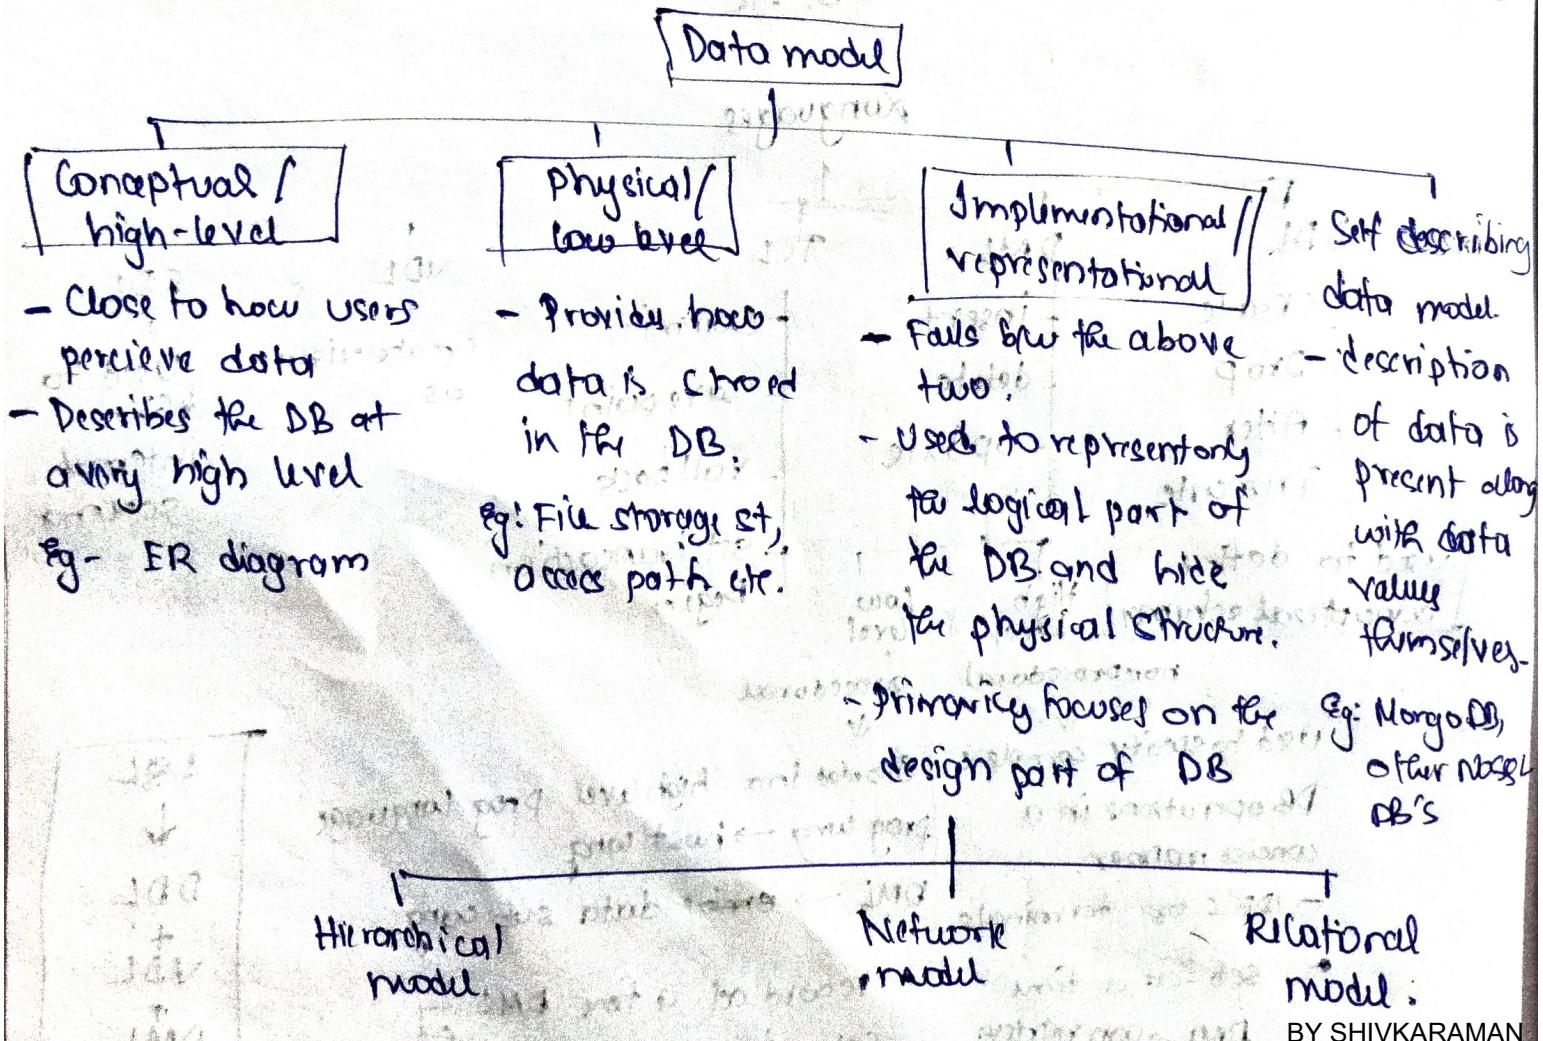
D BMS - Important

- * DBMS → Software to manage database.
 - * Functionality of DBMS → ~~Defining~~ Defining DB, creating DB, Manipulating DB and Sharing the DB.
 - * Limitations of file based approach
 - Data redundancy and inconsistency
 - Concurrent users not allowed
 - Multiple file formats
 - File path i.e. Data dependence.
 - * Characters of DBMS
 - Metadata and file structure stored in db catalog
 - Data is separate from program
 - Supports multiple views
 - Access controls
 - * DBA → Schema design, access control, selection of appropriate software to handle DB, etc

Languages

DDL	DML	TCL	VDL	SDL
Create Drop Alter Truncate	insert delete update	commit save point roll back set transaction begin	create view as	Used to define internal schema
Used to define conceptual schema	High Level	Low Level		
	nonprocedural	procedural		
Used to specify complex DB operations in a concise manner.		Embedded in a high level prog language prog lang → host lang		
- IDE's or terminals		DML → entire data sub lang		
• set-at-a-time		• record-at-a-time DML		
DML → can retrieve many records		Can retrieve only 1 record at a time		

- * instance / current state → Snapshot of DB at current time
- * Schema → Overall design / description of DB. (fields, attributes)
- * Sub-schema / view → Subset of schema
- * Schema diagram → Diagrammatic representation of schema.
- * Schema construct → A component of schema or object within schema. Eg: STUDENT
- * DB state → Empty, initial, current, valid state.
- * schema evolution → adding a new field to schema.
- Relational → means arranges data into tables containing rows/records/tuples and columns/attributes. Relation b/w tables is through foreign key.
- Data model → Collection of concepts that can be used to describe the structure of DB (data types, relations, constraints etc.)



Conceptual data model

• 3-Schema architecture

Goal → To separate user applications from physical database.
The user application is at the top-most level and the physical database is at bottom-most level.

1) Internal Level

- Has internal schema
- Describes the physical storage structure of DB
- Uses physical data model.

2) Conceptual level

- uses conceptual schema
- Represents structure of whole DB
- Uses ~~representational~~ data model.
- Hides phy storage structure and focuses on entities, relations etc.

3) External level

- has a no of views / external schema
- representational data model.

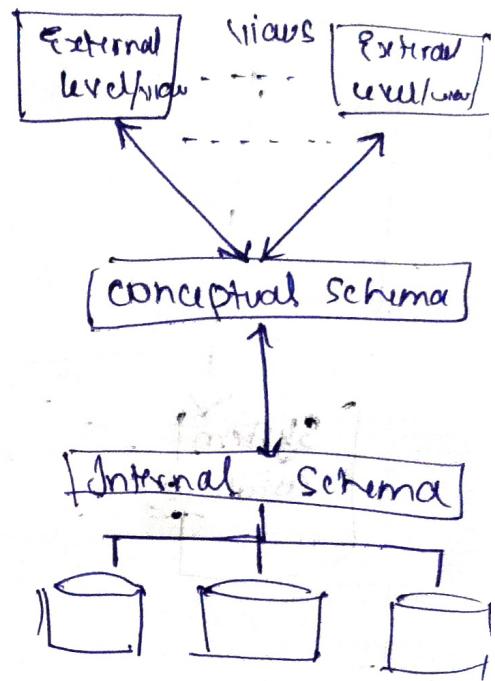
Transformation of request from external → conceptual → internal level
is called as mapping

• Data independence → Capacity to change Schema at one level of DBMS system without having to change schema at the next higher level.

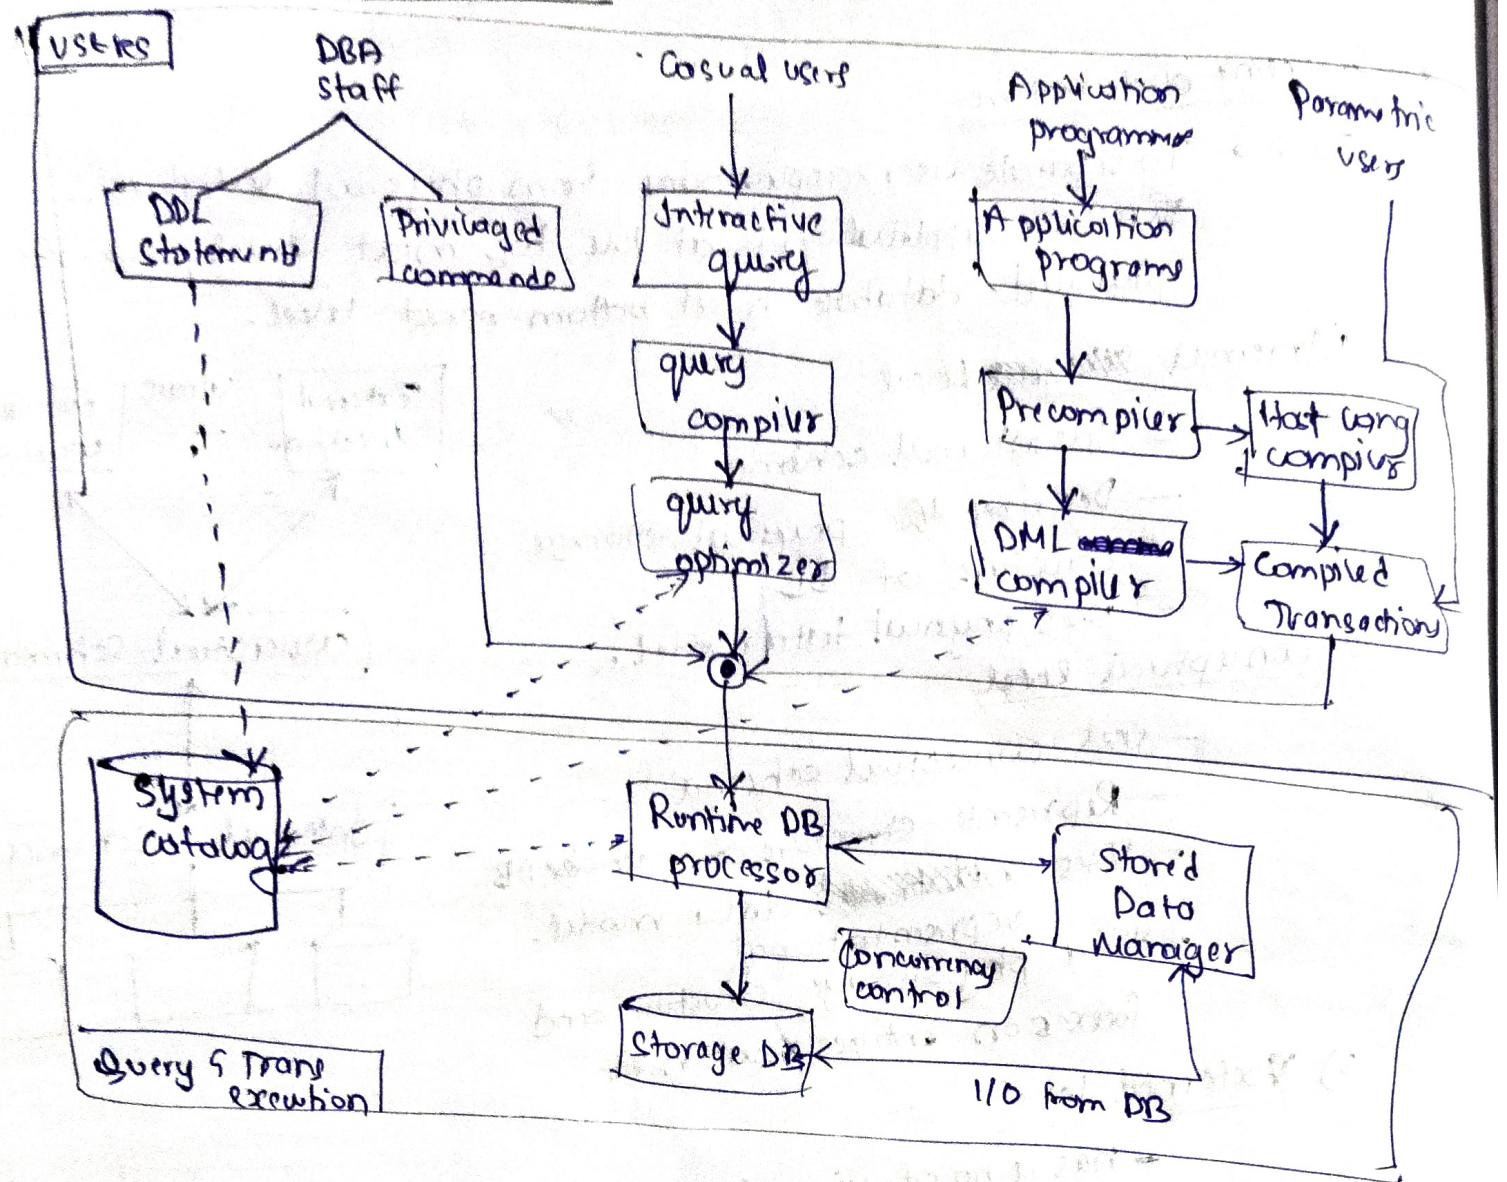
Logical data independence → Change in conceptual schema without change in external schema.

Physical data independence → Change in physical / internal schema without change in conceptual schema.

→ Harder to achieve



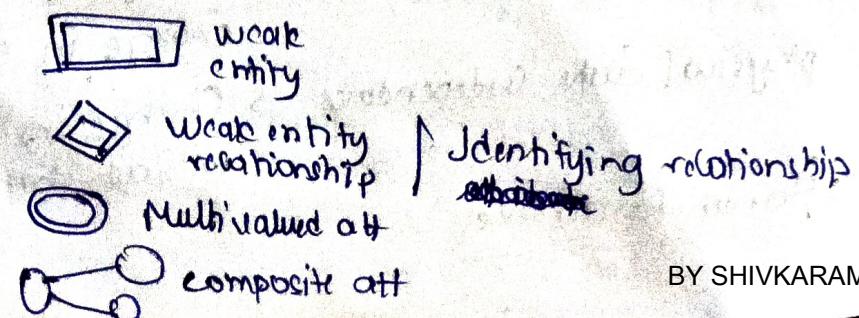
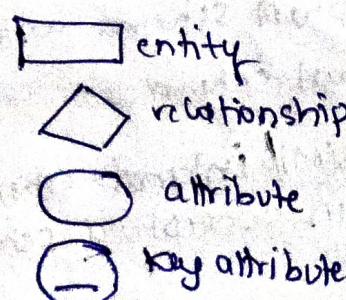
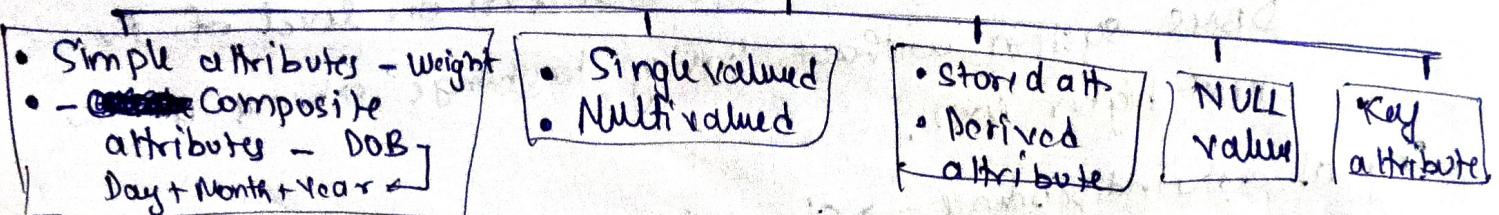
• DBMS component modules



ER diagrams

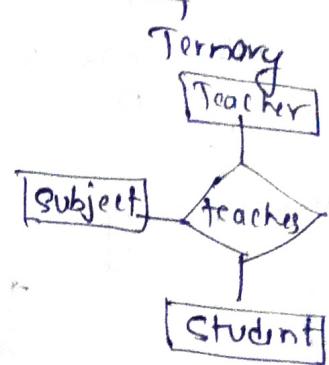
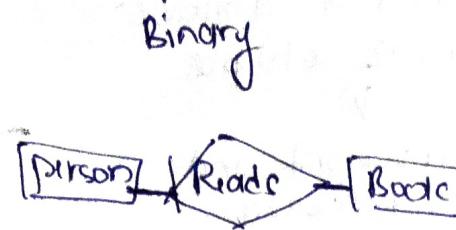
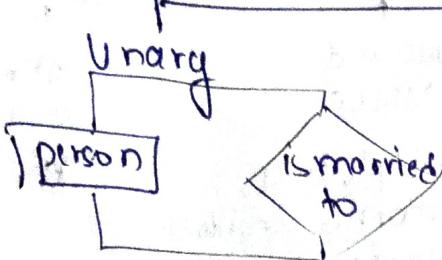
- Entity → A thing in real world
- Attributes → Properties of an entity.

Types of attributes



• relationship: association among 2 or more entities.

Degree of relationship



Relationship constraints

Cardinality Ratio

- 1:1
- 1:N
- N:1
- N:M

Participation constraint

Total

1

Partial



Keys

Candidate key

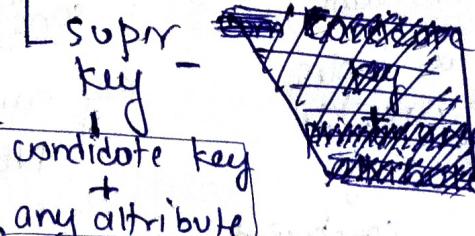
- Attribute whose value is unique and is capable of identifying a relationship \rightarrow Key
- Collection of all keys

Primary key

Most appropriate candidate key

Alternate keys
keys other than the PK

Foreign key.
Attribute or set of attributes that refers to PK of same table or different table.



Relational Model

rows \rightarrow tuples

cols \rightarrow attributes

table \rightarrow relation

domain \rightarrow datatype + range + format

Degree \rightarrow No of attributes

Cardinality \rightarrow No of tuples

Characteristics of a relation

Ordering of tuples inside a relation	Order of attribute inside a tuple	Values and NULLS	Interpretation of a relation
Not necessary	Important and necessary	<ul style="list-style-type: none"> - Each tuple in a relation is considered atomic → Multivalued att and composite att not allowed 	
	$T = \langle (\text{Roll No}, 2), (\text{Name}, \text{Bon}) \rangle$	<ul style="list-style-type: none"> Multi-val → separate rel val comp → Break down into simple components 	
	Not necessary when value along with attribute is given.	<ul style="list-style-type: none"> - NULL → Value is unknown or it does not apply 	

Constraints

Inherent model based constraints	Schema based / explicit constraint	Application based constraint
already present in data model. Eg: Duplicate tuples not allowed	Defined by the schema of data modul.	
Domain constraints data type + range + format	Key constraint No two tuples can have the same key.	Entity integrity constraint Key cannot be NULL
	Constraint Update operation and constraint violation	Referential Integrity constraint Tuple in one relation that refers to another relation must refer to an existing tuple in that relation. i.e
• Anomaly → Unexpected behavior during update/ delete/insert operations.		$FK = \text{NULL}$ or $FK = \text{value of PK of another relation}$

- 1) Insert → Domain, Key, Entity integrity, Referential integrity
- 2) Delete → Only referential integrity
- 3) Update → Referential integrity, Key constraint

Relational Algebra

Unary

- **SELECT (σ)**
Used to select columns
→ σ (Table) col-list

- **Project (Π)**
Used to select rows based on condition
→ Π (Table) condition

- **Rename (ρ)**

Binary Set

- **Union (\cup)**
Intersection (\cap)
Set difference ($-$)
Cartesian Product (\times)

- Two sets are union compatible if
 - Both have same no of attributes
 - Domain's are same
- Two sets must be union compatible to carry out \cup , \cap and $-$ operations

Binary

- **JOIN**

- EquiJoin
- Natural Join
- CrossJoin
- OuterJoin

RMS
condition

Other operation
Aggregate functions etc.

Division operation

Restrictions on updating views →

- Not allowed when view is created by joining more than table
- Not allowed on views defined with aggregate function
- Not allowed on views which don't include PK of base table.

SQL Injection

Based on $1=1$

User id: $105 \text{ OR } 1=1$

Password: $105 \text{ OR } 1=1$

SELECT * FROM users

WHERE user_id = 105 OR 1=1

SELECT * FROM password

WHERE password = 105 OR 1=1

1 is always true and will fetch all user id & password

Based on $"=1"$

User id: $"105" = 1$

Password: $"105" = 1$

SELECT * FROM users
WHERE user_id = "105" OR "1" = "
AND pass = "105" OR "1" = "

Will return all records of user_id and password

BY SHIVKARAMAN

Delete VS Truncate

DELETE

- DML command
- Deletes a row based on given condition.

**DELETE FROM table-name
WHERE condition;**

- Deleted rows are logged in MySQL transaction log, making it possible to rollback if needed.

TRUNCATE

- DDL command
- Used to delete all rows from table.

TRUNCATE TABLE

- TRUNCATE is not logged in transaction log which makes it unable to be rolled back.
- Cannot use WHERE clause.
- Faster than DELETE

Relational DB design

Design goals

Information preservation

Minimizing redundancy

Design Guidelines

Importing clear semantics to attributes.

- Don't combine attribute from multiple entity types in table

Inception anomaly

- Certain information cannot be stored until some other info is also stored

Redundant info in tuples & update anomalies.

- anomaly - Unusual behavior during DML Operations.

NULL values

- Problems
 - Wasted storage
 - Problem in understanding
 - Problem when joining operation
- Generation of spurious tuples
- Bad Join

Sol → Avoid fat tables

→ Design schema that doesn't suffer from any anomaly

Deletion anomaly

- Deletion of data representing certain information results in losing some other info as well.

Update anomaly

- presence of redundant data will lead to inconsistency

BY SHIVKARAMAN

- Functional dependency :- It is a constraint or relationship between two sets of attribute from DB.

$X \rightarrow Y$: X is functionally dependent on Y

Trivial FD: Y is a subset of X $\Rightarrow X \cap Y \neq \emptyset$

Non-trivial FD: $X \cap Y = \emptyset$

- Normalization

Process of analyzing the given schema based on their FD's and PK's to achieve desirable properties of minimizing redundancies and information preservation.

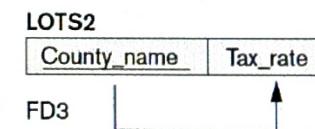
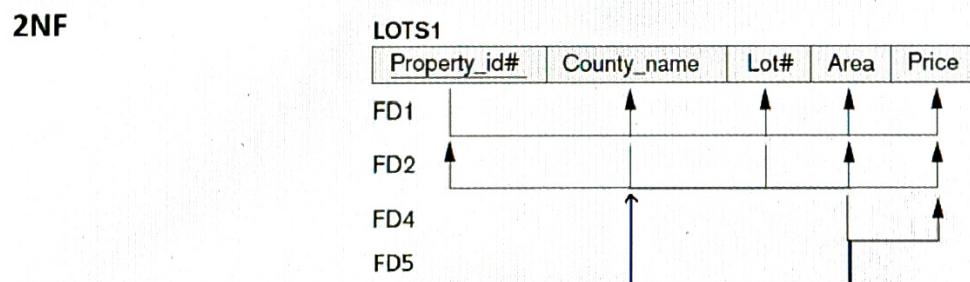
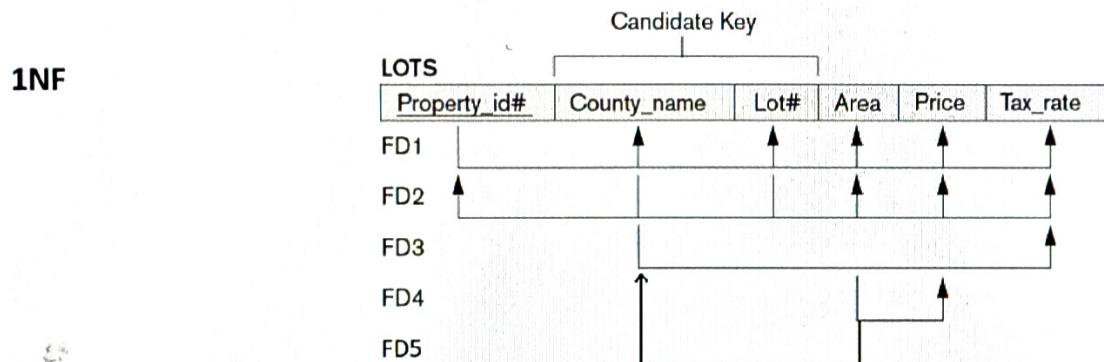
1NF: No multivalued attributes. i.e all attributes are atomic

2NF: No partial dependency i.e non-prime attribute not partially dependent on any of the key.

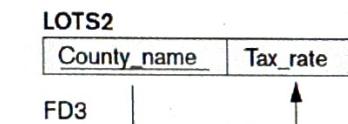
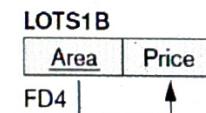
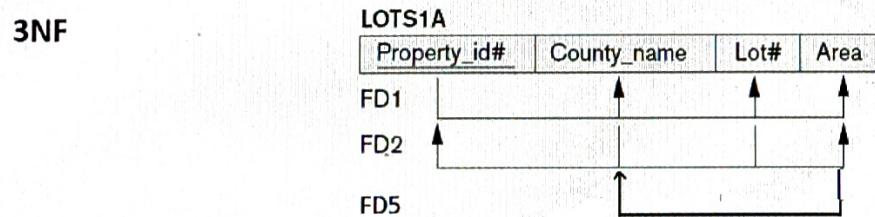
3NF: No transitive dependency i.e no nonprime attribute is transitively dependent on any key.

BCNF: Candidate key \rightarrow Non-prime attribute ~~the dependency should not be present~~

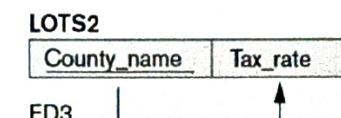
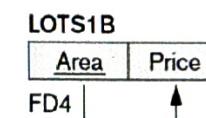
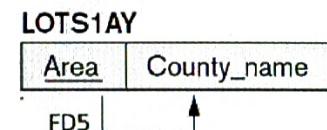
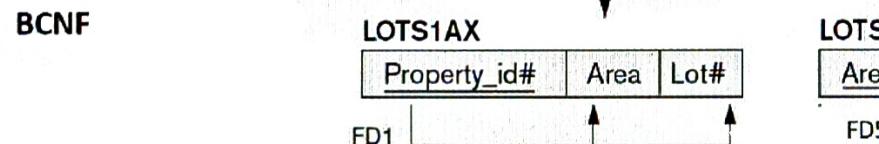
We can decompose LOTS1A into two BCNF relations LOTS1AX and LOTS1AY. This decomposition loses the functional dependency FD2 because its attributes no longer coexist in the same relation after decomposition



FD3: PARTIAL FUNCTIONAL DEPENDENCY
So we create a new table for the FD



FD4 is leading transitive FD => we create a new relation for it



FD5: A candidate key is FD on a non prime attributes

Serializability

When multiple transactions are being executed by the operating system in a multiprogramming environment, there are possibilities that instructions of one transaction are interleaved with some other transaction.

Schedule – A chronological execution sequence of a transaction is called a schedule. A schedule can have many transactions in it, each comprising of a number of instructions/tasks.

Serial Schedule – It is a schedule in which transactions are aligned in such a way that one transaction is executed first. When the first transaction completes its cycle, then the next transaction is executed. Transactions are ordered one after the other. This type of schedule is called a serial schedule, as transactions are executed in a serial manner.

In a multi-transaction environment, serial schedules are considered as a benchmark. The execution sequence of an instruction in a transaction cannot be changed, but two transactions can have their instructions executed in a random fashion. This execution does no harm if two transactions are mutually independent and working on different segments of data; but in case these two transactions are working on the same data, then the results may vary. This ever-varying result may bring the database to an inconsistent state.

To resolve this problem, we allow parallel execution of a transaction schedule, if its transactions are either serializable or have some equivalence relation among them. - **view equivalence and conflict equivalence**

1) Conflict Equivalence

Conflicts between transactions typically arise when they access shared data concurrently.

Read-Write, Write-Read, Write-Write \rightarrow Conflicts

Read-Read \rightarrow Non conflict Pair

Compare adj-operations \rightarrow If non-conflict \rightarrow Swap

conflicting \rightarrow NO SWAP

Eg: S Check if S and S' are conflict equivalent

T_1	T_2
R(A)	
W(A)	R(A)

$R(B)$

Non conflict
 \Rightarrow Swap

T_1	T_2
R(A)	
W(A)	R(A)

$R(B)$

Non conflict
 \Rightarrow Swap

$$\therefore S \equiv S'$$

T_1	T_2
R(A)	
W(A)	R(A)

$R(B)$

~~NON CONFLICT~~ = S'

T_1	T_2
R(A)	
W(A)	R(A)

$R(B)$

Concurrency Control

Lock Based Protocols

Time Stamp based protocols

- Binary locks: locked or unlocked
- Shared / exclusive locks: This type of locking mechanism differentiates the lock based on their uses.

Write operation → exclusive lock
Read operation → shared lock.

1) Simplistic lock protocol: lock before write operation and unlock after write operation is completed.

2) Precalming lock protocol: Pre claiming lock protocols evaluate ~~the sequence of operations of transaction~~ their operations and create a list of data items on which they need locks. Before initializing an execution, the transaction requests the system for all the locks it needs before hand. If all locks are granted, the transaction executes and releases all locks when all operations are over. If all locks are not granted, the transaction rolls back and waits until all locks are available.

3) 2-Phase locking: Divides the execution phase of a transaction into three parts

1) Growing phase: As transaction starts executing, it seeks permission for the lock it requires

2) Second part: Transaction acquires all the locks.

3) Shrinking phase: As soon as the transaction releases its first lock, third phase starts. In this phase, the transaction cannot demand any new lock, if only releases acquired locks.

→ Strict 2-P-locking: Same as 2-PL but here it does not release any lock until the commit point and releases at that time.

How data is stored in DBMS

1) Clustered Index Structure

For given table Employee [Emp-id | Name | Email | Dept]

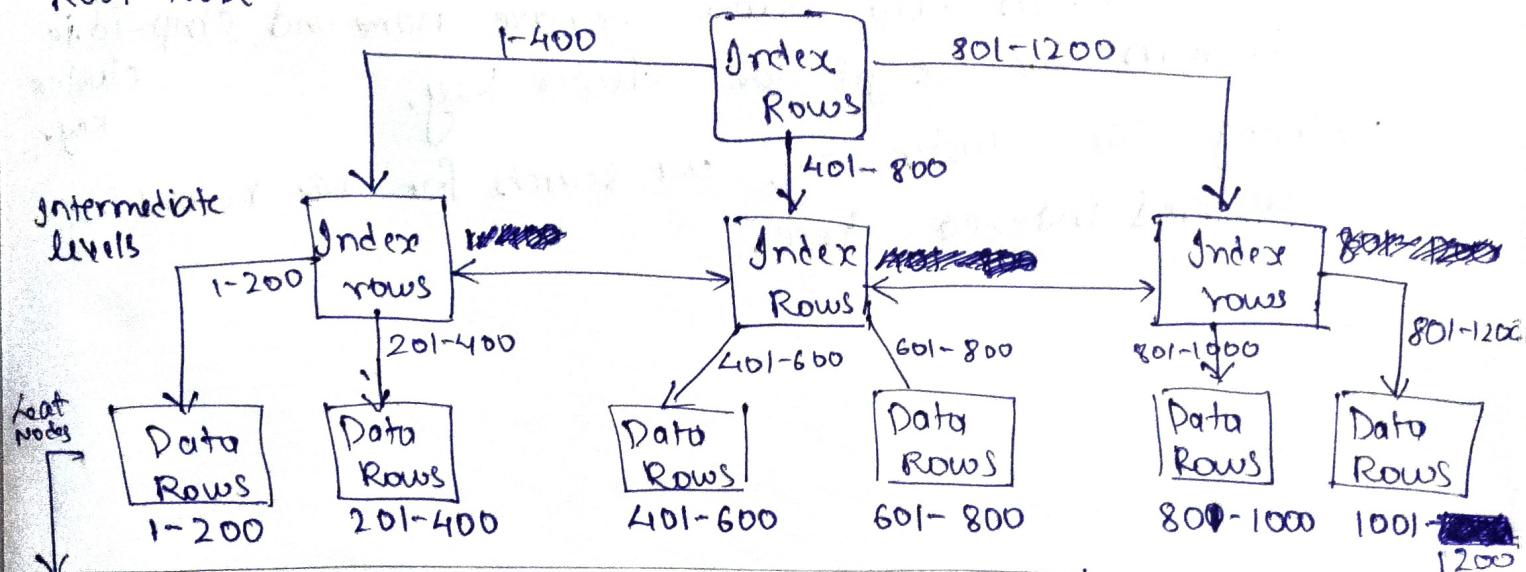
→ 1200 rows

* First, the tuples/rows are sorted according to the PK

* The data is stored in a ~~tree like~~ series of data pages in a tree like structure → Called as

B-tree / Index B-tree / Clustered structure

ROOT NODE



Data pages: 8 KB → Table data actually resides here!

* So in clustered index, indexing is done on the primary key. So this is helpful when we search ~~using~~ by primary key.

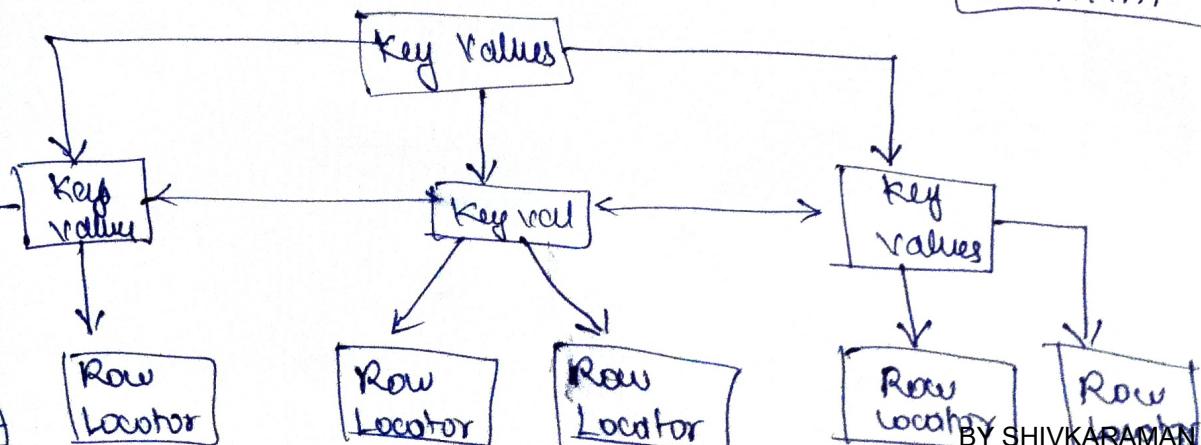
2) Non-Clustered index structure

* What if we search by Employee's name? There is no index for name column → SQL server has to scan every record in the table.

* Solution → form a non clustered index on name column

Inefficient

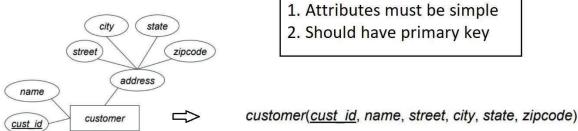
Name	Emp-id
Alexa	932
Becky	421
Charlie	765



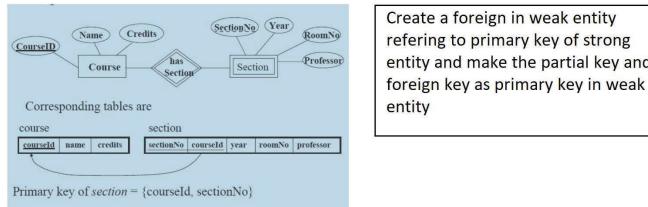
- * Here since we are searching on name column, the key values i.e. name is sorted in alphabetical order. The row locator present at the leaf nodes contain employee names and its ~~cluster~~ key i.e. Employee ID.
- * So now when we search employee by name, both non-clustered index and clustered index are used to fetch data.
 - 1) Use non-clustered index on name column to search for given name. In Row locator, we have name and emp-id i.e. So from this we get the cluster key.
 - 2) Using the cluster key, we search for the record in clustered index tree.

ER Model -> Relational Model

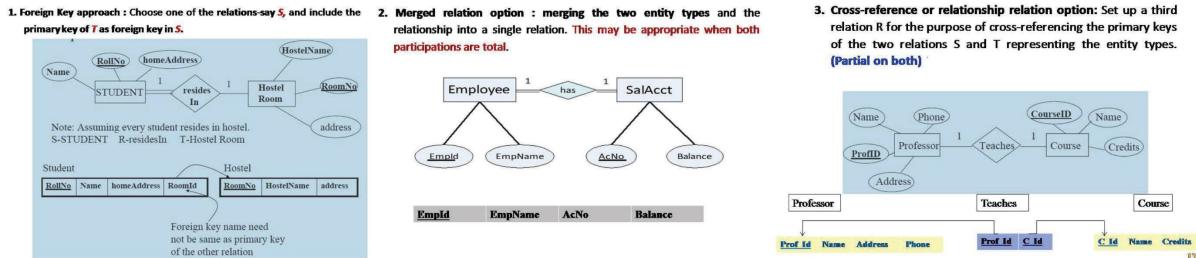
Step 1: Mapping of Regular Entity Types.



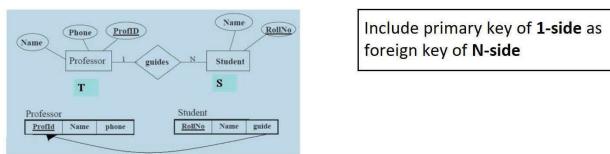
Step 2: Mapping of Weak Entity Types



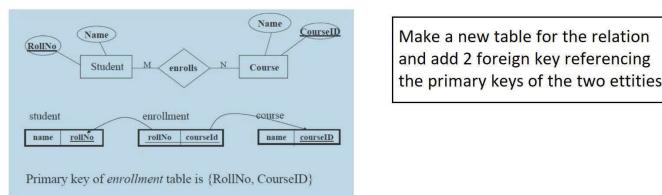
Step 3: Mapping of 1:1 Relation Types



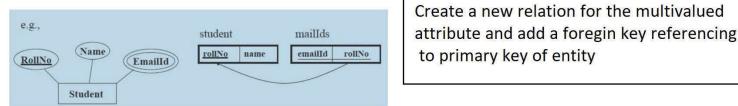
Step 4: Mapping of 1:N Relation Types



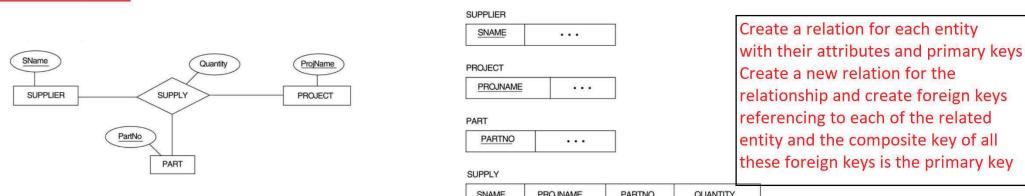
Step 5: Mapping of M:N Relationship Types.



Step 6: Mapping of Multivalued attributes.

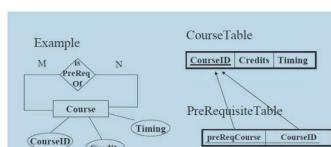


Step 7: Mapping of N-ary Relationship Types.



Handling Recursive relationships

- Make a table T for the participating entity set E and one table for recursive relationship R.



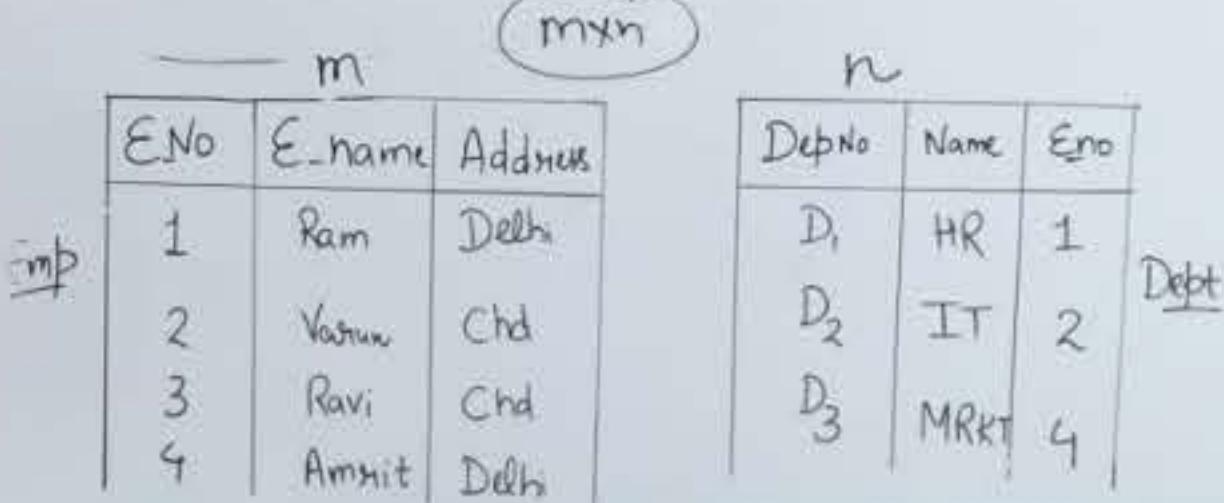
BY SHIVKARAMAN

JOIN = CROSS PRODUCT + SELECTION BASED ON CONDITION

Emp		Dept	
ENo	Ename	DepNo	Eno
1	Ram	D ₁	1 ✓
1	Ram	D ₂	2
1	Ram	D ₃	4 ✓
2	Vasun	D ₁	1
2	Vasun	D ₂	2 ✓
2	Vasun	D ₃	4
3	Ravi	D ₁	1
3	Ravi	D ₂	2
3	Ravi	D ₃	4 ✓
4	Amit	D ₁	1
4	Amit	D ₂	2
4	Amit	D ₃	4 ✓

'Natural JOIN'

"find the Emp Names who is working in a department"



A natural join is a type of join operation in relational database systems that combines rows from two or more tables based on columns with matching names and data types. Unlike other join types where you explicitly specify the join condition, in a natural join, the database system automatically matches columns with the same name in the joined tables.

SELECT Ename from Employee
NATURAL JOIN Dept;

Table joined with itself

Self JOIN

T ₁	T ₂
S ₁ C ₁	S ₁ C ₁
S ₁ C ₁	S ₂ C ₂
S ₁ C ₁	S ₁ C ₂ ✓
S ₂ C ₂	S ₁ C ₁
S ₂ C ₂	S ₂ C ₂
S ₂ C ₂	S ₁ C ₂
S ₁ C ₂ S ₁ C ₂	
S ₁ C ₂ S ₂ C ₂	

find Student id who is enrolled in at least two Courses.

Join = Cross product
+ Condition

Select

from Study as T₁,
Study as T₂]

F_{and} --> Where T₁.S_id = T₂.S_id
and T₁.C_id <> T₂.C_id

S_id	C_id	Since
S ₁	C ₁	2016
S ₂	C ₂	2017
S ₁	C ₂	2017

Alternate approach

```
SELECT sid
FROM study
GROUP BY sid
HAVING COUNT(cid) >= 2;
```

Emp		Dept	
ENo	Ename	Address	DepNo
1	Ram	Delhi	D ₁ , Delhi 1
1	Ram	Pune	D ₂ Pune 2
1	Ram	Patna	D ₃ Patna 4
2	Vasun	Chd	D ₁ , Delhi 1
2	Vasun	Pune	D ₂ Pune 2
2	Vasun	Patna	D ₃ Patna 4
3	Ravi	Chd	D ₁ , Delhi 1
3	Ravi	Pune	D ₂ Pune 2
3	Ravi	Patna	D ₃ Patna 4
4	Amit	Delhi	D ₁ , Delhi 1
4	Amit	Pune	D ₂ Pune 2
4	Amit	Patna	D ₃ Patna 4

Join = Cross Prod + Condition

Equi JOIN =

find the Emp name who worked in a department having location same as their address?

ENo	Ename	Address	DepNo	Location	Eno
1	Ram	Delhi	D ₁	Delhi	1
2	Vasun	Chd	D ₂	Pune	2
3	Ravi	Chd	D ₃	Patna	4
4	Amit	Delhi	D ₁	Delhi	1

Select Ename from Emp, Dept Where

Emp.E_no = Dept.E_no And
Emp.Address = Dept.Location,

An equi join, also known as inner join, is a type of join operation in relational database systems that combines rows from two or more tables based on matching values in specified columns.

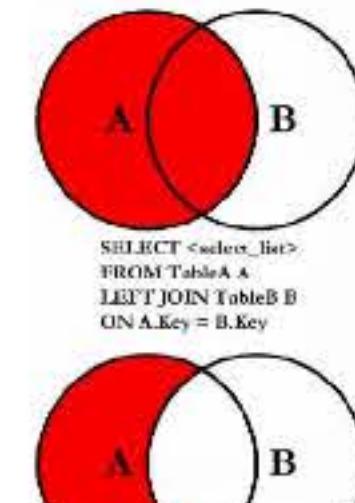
1. Left outer join

Keeps every tuple in first or left relation R, if no matching tuple is found in S, then attributes of S in the join result are filled or "padded1" with null values.

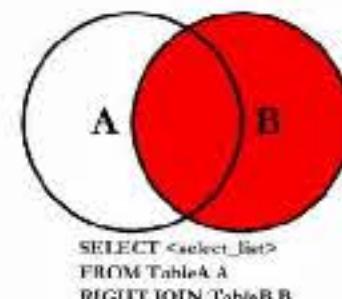
R	COL1	COL2
A	1	
B	2	
C	3	
D	4	
E	5	

S	COLA	COLB
A	1	
C	2	
D	3	
E	4	

R	COL1	COL2	COLA	COLB
A	1		1	
B	2		NULL	NULL
C	3		3	
D	4		4	
E	5		NULL	NULL



SQL JOINS



SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key

SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL

OR
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL

SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL

2. Right outer join

Keeps every tuple in second or right relation S in the result,

R	COL1	COL2
A	1	
B	2	
C	3	
D	4	
E	5	

S	COLA	COLB
A	1	
C	2	
D	3	
E	4	
F	5	

R	COL1	COL2	COLA	COLB
A	1		1	
B	2		NULL	NULL
C	3		3	
D	4		4	
E	5		NULL	NULL
F	6		6	

3. Full outer join

Keeps all tuples in both left and right relations when no matching tuples are found, padding them with null values as needed

R	COL1	COL2

<tbl_r cells="3" ix="2" max

What is Cursor in SQL ?

Cursor is a Temporary Memory or Temporary Work Station. It is Allocated by Database Server at the Time of Performing DML(Data Manipulation Language) operations on Table by User. Cursors are used to store Database Tables. There are 2 types of Cursors: Implicit Cursors, and Explicit Cursors. These are explained as following below.

1. Implicit Cursors:

Implicit Cursors are also known as Default Cursors of SQL SERVER. These Cursors are allocated by SQL SERVER when the user performs DML operations.

2. Explicit Cursors :

Explicit Cursors are Created by Users whenever the user requires them. Explicit Cursors are used for Fetching data from Table in Row-By-Row Manner.

SQL QUERIES

write an sql query to get the third maximum salary of an employee from employee table

Assuming you have an "Employee" table with columns "EmployeeID" and "Salary", you can use the following SQL query to get the third maximum salary:

```
sql  SELECT Salary
  FROM Employee
 ORDER BY Salary DESC
 LIMIT 1 OFFSET 2;
```

Explanation:

- The "ORDER BY Salary DESC" clause sorts the salaries in descending order.
- The "LIMIT 1 OFFSET 2" clause selects only the third row (salary) from the sorted result set. The "OFFSET 2" skips the first two rows (salaries), and the "LIMIT 1" selects the third row (salary).

How can you create an empty tabel from an existing table

You can create an empty table from an existing table in MySQL by using the "CREATE TABLE" statement with the "LIKE" keyword to copy the structure of the existing table. Here's an example:

```
sql  CREATE TABLE new_table LIKE existing_table;
```

This statement creates a new table called "new_table" with the same structure as "existing_table", but without any data.

You can also add additional table options to the "CREATE TABLE" statement, such as indexes, primary keys, and foreign keys, to further customize the new table.

Once the new table has been created, you can insert data into it using the "INSERT INTO" statement or by using other SQL commands to manipulate the data as needed.

How to fetch common records from two tables

To fetch common records from two tables, you can use the "INNER JOIN" clause in your SQL query. Here's an example:

Suppose you have two tables called "Table1" and "Table2" with a common column called "ID", and you want to fetch records that exist in both tables. You can use the following SQL query:

```
sql  SELECT Table1.*, Table2.*
  FROM Table1
 INNER JOIN Table2
    ON Table1.ID = Table2.ID;
```

How to fetch alternate records from a table

To fetch alternate records from a table, you can use the "ROW_NUMBER()" function along with a modulo operation in your SQL query to number the rows in the table, and then filter by even or odd row numbers. Here's an example:

Suppose you have a table called "Employee" with a column called "EmployeeID", and you want to fetch alternate rows from the table. You can use the following SQL query:

```
sql  SELECT *
  FROM (
    SELECT *, ROW_NUMBER() OVER (ORDER BY EmployeeID) AS rownum
    FROM Employee
  ) AS numbered_rows
 WHERE numbered_rows.rownum % 2 = 0;
```

Explanation:

- The "ROW_NUMBER()" function assigns a sequential number to each row in the table, based on the order specified in the "ORDER BY" clause.
- The "AS rownum" alias gives a name to the row number column.
- The "WHERE" clause filters the result set to only include rows with even row numbers, using the modulo operator "% 2" to check for divisibility by 2.

Note that you can customize the ordering of the rows by changing the "ORDER BY" clause.

How to fetch unique records from a table

To fetch unique records from a table, you can use the "DISTINCT" keyword in your SQL query

```
sql  SELECT DISTINCT CustomerID
  FROM Orders;
```

What is the command used to fetch first five characters of the string

To fetch the first five characters of a string, you can use the "LEFT" function in your SQL query. Here's an example:

```
sql  SELECT LEFT(ProductName, 5) AS ShortName
  FROM Products;
```

Explanation:

- The "LEFT" function returns the specified number of characters from the beginning of a string.